

Challenge 2025 – ADAMMA: On-Device MET Class Prediction from Smartphone Accelerometers

Arda Baris Basaran¹

¹ École Polytechnique Fédérale de Lausanne (EPFL)

ardabarisbasaran@hotmail.com

Abstract

We present a mobile system that continuously estimates a user’s Metabolic Equivalent of Task (MET) *class* (Sedentary, Light, Moderate, Vigorous) from smartphone accelerometer data and displays the cumulative time spent in each class throughout the day. The model is trained and evaluated on the WISDM “Activity Prediction” dataset, which contains six activities recorded at 20.00 Hz from phones carried in users’ front pockets [1, 2]. We map activities to MET classes using the Compendium of Physical Activities’ conventional intensity thresholds (sedentary 1.0–1.5, light 1.6–2.9, moderate 3.0–5.9, vigorous ≥ 6.0) [3, 4]. Our pipeline uses 5.00 s windows with a 1.00 s stride, a 43-dimensional handcrafted feature set implemented in PyTorch for fast on-device computation, and both classic ML classifiers (logistic regression, kNN, RBF-SVM, random forest, Gaussian NB, and a small MLP) and a compact 1D CNN suitable for on-device deployment. We report standard metrics (accuracy, precision, recall, F1) and per-sample inference latency, and describe an Android foreground-service implementation that maintains near real-time predictions and a clean UI with live cumulative time bars. We release source code (Android app, training code), trained models (including ONNX exports), and scripts to reproduce our results.

1 Introduction

Daily time spent across intensity classes (sedentary, light, moderate, vigorous) is a strong predictor of cardio-metabolic risk and all-cause mortality. METs (Metabolic Equivalents of Task) provide a simple, widely used scale to relate activity energy cost to resting metabolism. In the Compendium, intensities are typically categorized as sedentary 1.0–1.5 METs, light 1.6–2.9, moderate 3.0–5.9, and vigorous ≥ 6.0 [3, 4].

Modern smartphones ship with tri-axial accelerometers and sufficient compute to run lightweight ML models continuously. Our goal is a fully on-device app that (i) infers the current MET *class* in near real time from acceleration alone and (ii) maintains an always-

updating tally of cumulative minutes per class for “today”, with a small, readable UI and low battery overhead.

We build on the WISDM “Activity Prediction” dataset [1, 2]. While WISDM contains six activity labels—Walking, Jogging, Upstairs, Downstairs, Sitting, Standing—health guidance and the ADAMMA challenge operate in terms of *MET classes*. We therefore adopt a principled mapping from activities to intensity bins, train models to predict the four-class intensity label, and deploy a compact classifier on Android.

Contributions.

- A complete training & evaluation pipeline for four-class MET intensity recognition using WISDM accelerometer data.
- A 43-feature, PyTorch-implemented extractor adapted from the original WISDM feature set [1].
- Classic ML baselines and a compact 1D CNN, with accuracy/F1 and per-sample latency suitable for on-device use.
- An Android architecture (foreground service + live UI) for continuous sensing and cumulative time tracking.
- Reproducibility: scripts, configuration, and model artefacts.

2 Methodology

2.1 Dataset

We use the WISDM Activity Prediction dataset [2], collected with Android phones carried in the *front pants pocket* at 20.00 Hz (one sample every 50.00 ms). The canonical WISDM processing derives 43 features on 10 s windows (≈ 200 samples) [1], while our on-device pipeline (below) uses shorter windows for lower latency. The dataset provides six activities: *Walking, Jogging, Upstairs, Downstairs, Sitting, Standing*. According to the dataset metadata there are 1,098,207 labeled samples across 36 users (IDs 1 .. 36) with the following distribution:

Activity	Count	Percent
Walking	424,400	38.6%
Jogging	342,177	31.2%
Upstairs	122,869	11.2%
Downstairs	100,427	9.1%
Sitting	59,939	5.5%
Standing	48,395	4.4%

Each line in the WISDM dataset has the form:

```
[user], [activity], [timestamp],
[x-acceleration], [y-accel], [z-accel];
```

Fields:

- `user` — nominal, integers 1..36.
- `activity` — nominal in {Walking, Jogging, Sitting, Standing, Upstairs, Downstairs}.
- `timestamp` — numeric; generally phone uptime in nanoseconds (future datasets may use milliseconds since Unix epoch).
- `x-acceleration`, `y-accel`, `z-accel` — floating-point values typically in $[-20, 20]$. A value of 10 corresponds to $1g \approx 9.81 \text{ m/s}^2$; values include gravitational acceleration, so a resting vertical axis registers around ± 10 .

Dataset choice. Before committing to WISDM, we considered several other widely used activity recognition datasets. The UCI-HAR dataset [5] provides accelerometer and gyroscope signals from a waist-mounted smartphone, pre-segmented into 2.56 s windows; however, it lacks vigorous-intensity classes (no jogging or running). The MHEALTH dataset [6] includes rich multi-sensor streams (ECG, accelerometer, gyroscope, magnetometer) and diverse activities such as cycling and running, but its setup with multiple body-worn devices makes it less representative of a single-phone scenario. PAMAP2 [7] similarly offers detailed multi-sensor data and a broad range of activities mapped to METs, yet requires three body-mounted IMUs. More recently, WEEE [8] combines accelerometry with indirect calorimetry for precise energy expenditure, but is not phone-based. Compared to these, WISDM [1, 2] is uniquely well aligned with our goals: it uses only the smartphone’s internal accelerometer, sampled at 20 Hz, in the front pocket carry position. Since our aim is a deployable Android app using just the phone’s own sensors, WISDM provides the most realistic match in terms of modality and placement, while still allowing a principled mapping from activities to MET intensity classes (Table 1).

Why classification (not regression). We formulate intensity recognition as a *four-class* classification task rather than direct MET regression for several reasons.

First, the selected dataset (WISDM) does not provide continuous, ground-truth energy expenditure labels; it contains activity categories only, so any MET value would have to be *imputed* from external tables (e.g., the Compendium) and then treated as truth, introducing label noise and cascading bias [3, 4, 2]. Second, mapping a regressed MET *value* to intensity requires a second thresholding step (e.g., 1.5/3/6 METs); this two-stage pipeline compounds errors (regression + thresholding), whereas a single classifier can optimize the decision boundaries for the end objective directly. Third, absolute MET varies substantially with person-specific physiology (mass, fitness), biomechanics, and phone placement/orientation—factors that an accelerometer-only model cannot reliably disambiguate without per-user calibration. Treating the problem as classification avoids conveying spurious numerical precision and aligns the model’s outputs with clinical/public-health guidance and our UI.

Our windowing. For low-latency on-device prediction, we use **5 s** windows with a **1 s** stride at 20.00 Hz ($5 \times 20 = 100$ samples per window). This choice, aligned with typical mobile phone sensing setups, balances decision latency with sufficient temporal context to capture periodic structure in gait and stair-climbing, while providing more frequent predictions.

Preprocessing. We drop malformed or all-zero rows, cast types, and sort by `user`, `timestamp`. Windows are generated sequentially per user with a fixed stride (no overlap unless otherwise noted). Each window’s label is the *majority* activity over its samples. (Implementation in `datasets/dataset.py`.)

2.2 Mapping Activities to MET Classes

We map WISDM activities to intensity classes using Compendium thresholds [3, 4] and typical MET values:

Activity	MET Class (rationale)
Sitting, Standing	Sedentary (< 1.5 MET)
Walking (casual)	Light (1.6–2.9 MET)
Upstairs, Downstairs	Moderate (3.0–5.9 MET; stair climbing ~ 4 –8)
Jogging/Running	Vigorous (≥ 6.0 MET)

Table 1: Activity categories and their MET classes.

This yields four balanced-by-design intensity labels for the challenge: *Sedentary*, *Light*, *Moderate*, *Vigorous*.

2.3 Feature Extraction (43 dims)

We adopt the classic WISDM feature family [1], implemented in PyTorch (`models/model.py`) so the same code can run on CPU/GPU and be ported to mobile if

desired. For each window (shape $T \times 3$; $T=60$ with our 3 s windows) and each axis x, y, z :

- Mean, standard deviation, and average absolute deviation ($3 \times 3 = 9$).
- Average resultant acceleration $\mathbb{E}[\sqrt{x^2 + y^2 + z^2}]$ (1).
- Time-between-peaks per axis (3): local-peak detection; average spacing converted to ms via sampling rate.
- Binned distributions per axis (10 bins $\times 3 = 30$).

Total $9 + 1 + 3 + 30 = 43$ features per window. The extractor is window-length agnostic.

2.4 Models

We evaluate classic scikit-learn classifiers and a compact neural network suitable for mobile deployment:

- **Logistic Regression** (max_iter=1000).
- **kNN** ($k = 5$).
- **RBF-SVM**.
- **Random Forest** (100 trees).
- **Gaussian Naive Bayes**.
- **MLP** (hidden layers $128 \rightarrow 64$, ReLU, Adam, max_iter=200).
- **SimpleCNN1D** (ours): three depthwise-separable 1D convolution blocks with dilations 1/2/4 and ReLU6, global average pooling, and a small dense head ($64 \rightarrow 32 \rightarrow 4$). Trained with cross-entropy using Adam (lr 10^{-3}), dropout 0.1, for 10 epochs.

Training protocol. By default we use stratified 5-fold cross-validation *over windows* (Scikit’s StratifiedKFold, random_state=42). For fold k , we train on 4 folds and test on the held-out fold. The same protocol is applied to both ML models and the CNN. *Note:* this is not subject-exclusive; we discuss the implications in Section 7.

Metrics and latency. We report accuracy, weighted precision/recall/F1, and *per-sample* inference time (mean over the test set: batch predict \Rightarrow wall time / #samples, expressed in ms/sample). Code in `main.py`. For the CNN, we measure forward-pass wall time in evaluation mode.

2.5 Device Specifications

Build environment. The Android app is developed in Android Studio Narwhal 3 Feature Drop 2025.1.3, Android Gradle Plugin (AGP) 8.5.2, Gradle 8.7, with `compileSdk=34`, `targetSdk=34`, and `minSdk=24`.

Test device. We tested end-to-end sensing and the UI on a Samsung Galaxy A15. The app requests a 20.00 Hz accelerometer sampling rate; on this hardware, the sensor reports the closest supported rate, approximately

19.00 Hz. The window size and stride are set to **5 s** and **1 s**, respectively, matching the training setup.

3 App Design

The Android app follows a simple, readable flow:

- **Splash Screen:** Lightweight splash while the app initializes services and loads model assets.
- **Starter Screen:** Shows the app logo and a prominent *Start* button to begin sensing.
- **Activity Monitor:** The main screen presents the *current state* (Sedentary/Light/Moderate/Vigorous) with an icon and label. A compact *Debug Information* panel displays live accelerometer values (x, y, z), the measured sampling rate in Hz, the model’s output probability distribution over the four classes, and the window/stride configuration in seconds. These values update in real time as new sensor data arrives.
- **History Screen:** A four-tab view (*Day/Week/Month/Year*) summarizing time spent per intensity. The Day tab shows, for each class, a 24-segment bar (one per hour) and a *pie chart* of today’s distribution, plus a *calories today (net)* estimate computed from per-second labels and user profile (age, weight, height, sex)—see Section 3.1. Week/Month/Year aggregate into 7/30/12 segments respectively. A *Reset Data* action in Settings clears all history.

3.1 Calorie Estimation

What a MET means. A Metabolic Equivalent of Task (MET) is a unit that scales energy cost relative to quiet rest. By convention,

$$1 \text{ MET} \triangleq 3.5 \text{ ml O}_2 \text{ kg}^{-1} \text{ min}^{-1}$$

and gross caloric expenditure (including resting) can be estimated from METs and body mass.¹ Importantly, MET \rightarrow kcal depends primarily on *body mass*; height/age/sex mainly matter through basal metabolic rate (BMR), which is separate from activity energy.

From class posteriors to a per-second MET. Our on-device model outputs, each second t , a posterior over four MET classes

$$\mathbf{p}_t = [p_t^{(\text{sed})}, p_t^{(\text{light})}, p_t^{(\text{mod})}, p_t^{(\text{vig})}], \quad \sum_c p_t^{(c)} = 1.$$

We assign each class a representative MET intensity m_c based on the Compendium ranges and our activity mapping (*sedentary*: sitting/standing; *light*: walking; *moderate*: stairs; *vigorous*: jogging). We use mid-range values for stability:

$$m_{\text{sed}} = 1.2, \quad m_{\text{light}} = 2.5, \quad m_{\text{mod}} = 4.5, \quad m_{\text{vig}} = 8.0.$$

¹The factor 3.5 ml/kg/min and the caloric equivalent 1 L O₂ \approx 5 kcal lead to the 200 constant below.

The per-second *expected* MET is the posterior mean:

$$\widehat{M}_t = \sum_{c \in \{\text{sed}, \text{light}, \text{mod}, \text{vig}\}} p_t^{(c)} m_c.$$

To reduce jitter while remaining responsive, we apply an exponential moving average (EMA) with smoothing parameter $\alpha \in (0, 1]$:

$$\widetilde{M}_t = \alpha \widehat{M}_t + (1 - \alpha) \widetilde{M}_{t-1}, \quad \widetilde{M}_0 = \widehat{M}_0.$$

Gross vs. net calories (personalized). Let W be body mass (kg), H be height (cm), A age (years), and S a sex indicator (+5 for male, -161 for female). We first compute resting metabolic rate (RMR) using the Mifflin–St Jeor equation:

$$\text{RMR} = 10W + 6.25H - 5A + S \quad [\text{kcal/day}].$$

Converting to seconds:

$$\dot{E}_{\text{rest}} = \frac{\text{RMR}}{24 \cdot 3600} \quad [\text{kcal/s}].$$

This defines a personalized baseline: $1 \text{ MET} = \dot{E}_{\text{rest}}$.

Thus, the instantaneous gross and net caloric rates are

$$\dot{E}_t^{\text{gross}} = \widetilde{M}_t \cdot \dot{E}_{\text{rest}}, \quad \dot{E}_t^{\text{net}} = \max(\widetilde{M}_t - 1, 0) \cdot \dot{E}_{\text{rest}}.$$

Daily totals are obtained by summation:

$$E_{\text{day}}^{(\cdot)} = \sum_{t \in \text{day}} \dot{E}_t^{(\cdot)}.$$

Units sanity check (worked example). Example: $W = 70$ kg, $H = 175$ cm, $A = 25$ years, male ($S = +5$). Then

$$\text{RMR} = 10 \cdot 70 + 6.25 \cdot 175 - 5 \cdot 25 + 5 = 1673.75 \text{ kcal/day},$$

$$\dot{E}_{\text{rest}} = \frac{1673.75}{86400} \approx 0.0194 \text{ kcal/s}.$$

For $\widetilde{M}_t = 8$ (vigorous jog):

$$\dot{E}_t^{\text{gross}} = 8 \cdot 0.0194 \approx 0.155 \text{ kcal/s} \quad (\approx 9.3 \text{ kcal/min}),$$

which closely matches standard tables for jogging intensity.

4 Results

We evaluate the six classic classifiers *and* the compact 1D CNN under identical 5-fold CV. Unless noted, handcrafted features (43-D) with window length 5.00 s at 20.00 Hz are used. We report the mean and standard deviation over folds.

Model	Acc (%)	Prec (%)	Rec (%)	F1 (%)	Inf. time (ms)
Logistic Regression	85.4 ± 0.2	84.9 ± 0.2	85.4 ± 0.2	84.9 ± 0.2	0.0003
kNN (k=5)	76.2 ± 0.4	75.8 ± 0.3	76.2 ± 0.4	75.4 ± 0.4	0.0119
RBF-SVM	75.1 ± 0.2	75.4 ± 0.4	75.1 ± 0.2	69.2 ± 0.3	1.8195
Random Forest	98.0 ± 0.1	98.0 ± 0.1	98.0 ± 0.1	98.0 ± 0.1	0.0096
Gaussian NB	84.0 ± 0.3	84.0 ± 0.3	84.0 ± 0.3	83.9 ± 0.3	0.0006
MLP (128,64)	96.8 ± 0.5	96.9 ± 0.5	96.8 ± 0.5	96.8 ± 0.5	0.0015
SimpleCNN1D	90.6 ± 1.2	91.5 ± 1.0	90.6 ± 1.2	90.8 ± 1.1	1.8067

Table 2: 5-fold CV performance using handcrafted 43-D features (5 s windows, 1 s stride). Metrics are mean \pm std across folds (in %), inference time reported as mean \pm std (ms/sample).

Model	Acc (%)	Prec (%)	Rec (%)	F1 (%)	Inf. time (ms)
Logistic Regression	56.7 ± 0.1	60.2 ± 0.5	56.7 ± 0.1	53.5 ± 0.2	0.001
kNN (k=5)	72.2 ± 0.3	81.0 ± 0.3	72.2 ± 0.3	72.3 ± 0.3	0.057
RBF-SVM	93.5 ± 0.4	93.5 ± 0.4	93.5 ± 0.4	93.4 ± 0.4	5.223
Random Forest	88.5 ± 0.2	89.6 ± 0.2	88.5 ± 0.2	87.3 ± 0.3	0.015
Gaussian NB	62.3 ± 0.3	68.2 ± 0.3	62.3 ± 0.3	63.0 ± 0.2	0.003
MLP (128,64)	94.7 ± 0.3	94.7 ± 0.3	94.7 ± 0.3	94.7 ± 0.3	0.002
SimpleCNN1D	99.0 ± 0.2	99.0 ± 0.2	99.0 ± 0.2	99.0 ± 0.2	0.1702

Table 3: 5-fold CV performance using raw features (5 s windows, 1 s stride). Metrics are mean \pm std across folds (in %), inference time reported as mean \pm std (ms/sample).

Latency. For each model, we measure test-time mean inference latency (ms/sample) using batch prediction, then divide wall time by the number of test windows (see `main.py`). For the CNN, latency is measured with the network in evaluation mode (BatchNorm/Dropout frozen). Classic scikit-learn classifiers are timed on CPU, whereas the CNN is timed on GPU; these are offline desktop measurements and *not* on-device estimates.

5 Future Work

There are several directions in which this work can be extended to improve robustness, accuracy, and generalizability. These include signal preprocessing for orientation invariance, additional lightweight features, temporal smoothing strategies, cross-dataset evaluation, broader model architectures, and deployment-focused compression.

Pre-filtering for orientation and gravity robustness.

Future systems can incorporate a lightweight gravity compensation step to reduce sensitivity to phone orientation. An IIR low-pass estimate of gravity per axis (with cutoff ≈ 0.25 Hz) can be maintained, after which linear acceleration is computed as $\mathbf{a}_{\text{lin}} = \mathbf{a}_t - \mathbf{g}_t$, with the recursive update $\mathbf{g}_t = \alpha \mathbf{g}_{t-1} + (1 - \alpha) \mathbf{a}_t$, where $\alpha = \exp(-2\pi f_c / f_s)$. In addition, the magnitude $m = \|\mathbf{a}_{\text{lin}}\| = \sqrt{x^2 + y^2 + z^2}$ can be computed to obtain orientation-invariant cues.

Feature set extensions from linear acceleration.

The feature set can be enriched with additional time-domain, orientation-robust statistics computed per window from the signals x , y , z , and m . Candidate fea-

tures include mean, standard deviation, and median absolute deviation (16 features across four signals), interquartile range ($p_{75} - p_{25}$) for each signal (4 features), zero-crossing rate (4 features), signal magnitude area $SMA = \frac{1}{N} \sum_t (|x| + |y| + |z|)$ (1 feature), energy values $\text{mean}(x^2)$, $\text{mean}(y^2)$, $\text{mean}(z^2)$, $\text{mean}(m^2)$ (4 features), and jerk-based statistics such as the standard deviation of first differences Δa for each signal (4 features). These features are computationally efficient and can complement or replace raw input windows.

Temporal smoothing for stable MET classes. Because per-window predictions may fluctuate, class probabilities can be smoothed directly using an exponential moving average (EMA), $\hat{\mathbf{p}}_t = \alpha \mathbf{p}_t + (1 - \alpha) \hat{\mathbf{p}}_{t-1}$, with $\alpha \approx 0.7$. This can be combined with a simple hysteresis mechanism that requires k consecutive frames (e.g., $k = 2-3$) before committing to a new class. Together, these methods can yield steadier on-device predictions and more stable daily summaries, reducing jitter without hiding systematic errors.

Cross-dataset evaluation. Generalization can be evaluated by extending experiments beyond WISDM to additional benchmarks such as *UCI-HAR*, *MHEALTH*, *PAMAP2*, and *WEEE*. Sampling rates and windowing schemes can be harmonized, and activities can be mapped to MET classes for consistency. Where user identifiers are available, subject-exclusive splits can be used to quantify robustness under user-level distribution shift. Cross-dataset transfer can also be studied by training on one dataset and testing on another.

Model architectures beyond a simple 1D CNN. The current work is limited to basic classifiers and a simple 1D CNN, but many efficient time-series models can be explored in the future. Recurrent models such as LSTM and GRU with sequence-to-label pooling or attention can capture longer dependencies. Temporal Convolutional Networks (TCNs) with dilated causal convolutions can model long receptive fields. CNN-RNN hybrids such as DeepConvLSTM can combine convolutional feature extractors with recurrent backends. Multi-scale CNNs, including InceptionTime, ResNet1D, and FCN-style architectures with parallel kernels and residual connections, can provide richer representations. Transformer-based encoders, including lightweight or linear-attention variants such as Performer and Informer, can also be adapted for efficient on-device inference. Finally, self-supervised pretraining approaches (e.g., contrastive learning methods such as SimCLR) on large unlabeled motion datasets can provide strong feature extractors that require only a small classifier head for fine-tuning.

6 Software & Reproducibility

Repository layout.

- `app/` and `gradle/wrapper/`: Kotlin app with foreground service, `SensorManager`, feature computation, and UI.
- `METTrackerTrainer/`: contains the scripts to train and evaluate different ML/DL models `main.py`, `datasets/`, `models/`.
- `video-onnx-and-apk-report/`: contains the APK, a video demonstration of the app, the ONNX file `cnn.onnx`, and the report in PDF.

Running training. Example (handcrafted features, classic ML):

```
python main.py \
  --dataset WISDM \
  --dataset_path /path/to/WISDM/ \
  --feature_extraction handcrafted \
  --model ml \
  --n_splits 5 \
```

Run the compact CNN by switching `--model cnn`; use either handcrafted features or raw windows:

```
python main.py \
  --dataset WISDM \
  --dataset_path /path/to/WISDM/ \
  --feature_extraction raw \
  --model cnn \
  --n_splits 5 \
  --batch_size 32
```

7 Discussion & Conclusion

This work demonstrates the feasibility of a fully on-device system for recognizing daily physical activity intensities from smartphone accelerometers. Using the WISDM dataset, we mapped six activity labels into four clinically relevant MET classes (Sedentary, Light, Moderate, Vigorous), extracted a 43-dimensional handcrafted feature set, and compared several lightweight machine learning classifiers against a compact 1D CNN trained directly on raw windows.

Our experiments show that classic models such as random forests and MLPs achieve very high accuracy (up to $\sim 98\%$), with millisecond-level inference latency suitable for deployment. The compact CNN achieves strong accuracy on raw signals ($\sim 99\%$) while remaining efficient enough for near real-time use. These results highlight that both handcrafted features with simple models and compact deep networks are viable pathways for energy-efficient mobile sensing.

On-device evaluation on a mid-range Android phone confirmed that the full pipeline—sensor acquisition, feature extraction, inference, and UI updates—can run continuously with minimal overhead. The app provides

interpretable outputs (current class, posterior probabilities) and accumulates daily minutes per class, aligning directly with public health guidelines. A personalized calorie estimator extends utility by translating predicted METs into gross and net energy expenditure.

Nevertheless, several limitations remain. First, the dataset is restricted to a single phone placement (front pocket) and does not cover the full diversity of free-living activities. Second, our evaluation uses window-level cross-validation rather than subject-exclusive splits, so true generalization across users is not yet quantified. Third, the system currently assumes fixed class–MET mappings and does not adapt to individual differences in physiology or phone orientation.

Future directions include exploring orientation-invariant preprocessing, additional lightweight features, temporal smoothing strategies, subject-exclusive and cross-dataset validation, and more advanced but efficient sequence models (e.g., TCNs, DeepConvLSTM, lightweight transformers). These steps would strengthen robustness and ecological validity, paving the way toward real-world deployment at scale.

References

- [1] J. R. Kwapisz, G. M. Weiss, and S. A. Moore, “Activity recognition using cell phone accelerometers,” in *Proceedings of the Fourth International Workshop on Knowledge Discovery from Sensor Data (SensorKDD) at KDD-10*, (Washington, DC), 2010.
- [2] “Wisdm lab: Dataset (activity prediction).” <https://www.cis.fordham.edu/wisdm/dataset.php>. Accessed 2025-09-04.
- [3] S. D. Herrmann *et al.*, “2024 adult compendium of physical activities: A third update of activity codes and met values,” *Journal of Sport and Health Science*, 2024. Defines sedentary 1.0–1.5 METs, light 1.6–2.9, moderate 3.0–5.9, vigorous ≥ 6.0 .
- [4] B. E. Ainsworth, W. L. Haskell, S. D. Herrmann, *et al.*, “2011 compendium of physical activities: a second update of codes and met values,” *Medicine & Science in Sports & Exercise*, vol. 43, no. 8, pp. 1575–1581, 2011.
- [5] D. Anguita, A. Ghio, L. Oneto, X. Parra, and J. L. Reyes-Ortiz, “Smartphone-based recognition of human activities and postural transitions,” *Pattern Recognition Letters*, vol. 34, no. 16, pp. 2183–2193, 2013.
- [6] O. Banos, R. Garcia, J. A. Holgado-Terriza, M. Damas, H. Pomares, and I. Rojas, “mhealth-droid: a novel framework for agile development of mobile health applications,” in *International Workshop on Ambient Assisted Living*, pp. 91–98, Springer, 2014.
- [7] A. Reiss and D. Stricker, “Benchmarking real-world activity recognition for mobile devices,” in *2012 16th International Symposium on Wearable Computers*, pp. 1–8, IEEE, 2012.
- [8] K. Ailneni and G. Beltrame, “The wearable energy expenditure (weee) dataset: benchmarking energy expenditure estimation from wearable sensors,” *Scientific Data*, vol. 10, no. 1, p. 95, 2023.